

HAN LIN

✉ hanlincs@cs.unc.edu

🏠 hl-hanlin.github.io

🎓 citation:1652 | h-index:13

INTERESTS

Generative models, World models, Multimodal learning, Representation Learning, Reinforcement Learning, Theory-grounded algorithms for efficient Transformers

EDUCATION

University of North Carolina at Chapel Hill

Ph.D. in Computer Science

2023 - Exp. 2028

GPA: 4.0/4.0

- MURGe-Lab. Advised by Prof. [Mohit Bansal](#)

Columbia University

M.S. in Computer Science

2021 - 2023

GPA: 4.13/4.33

- DVMM Lab. Advised by Prof. [Shih-Fu Chang](#)
- ROAM Lab. Advised by Prof. [Matei Ciocarlie](#) and Prof. [Shuran Song](#)

Relevant Courses: Learning Theory, Algorithms, Machine Learning, Unsupervised Learning, Bandits & Reinforcement Learning, Causal Inference, Computer Vision, Robotics Learning

Columbia University

M.S. in Financial Engineering

2018 - 2020

GPA: 4.04/4.33

- Advised by Prof. [Krzysztof Choromanski](#) from Google Deepmind

Relevant Courses: Optimization, Combinatorial Optimization, Stochastic Models, Stochastic Calculus, Monte Carlo Methods, Statistical Inference, Bayesian Statistics, Graphical Models

Central University of Finance and Economics

B.S. in Financial Engineering

2014 - 2018

GPA: 90.43/100

Relevant Courses: Linear Algebra, Mathematical Analysis, Probability, Statistics, Real Analysis, Numerical Methods, Stochastic Process, Differential Equations

PREPRINTS

1. Seeing Isn't Knowing: Do VLMs Know When Not to Answer Spatial Questions (and Why)?

[\[arxiv\]](#)[\[project page\]](#)

Yue Zhang, Zun Wang, **Han Lin**, Jialu Li, Yonatan Bitton, Idan Szepkter, Mohit Bansal

2. Planning with Sketch-Guided Verification for Physics-Aware Video Generation

[\[arxiv\]](#)[\[github\]](#)

Yidong Huang, Zun Wang, **Han Lin**, Dong-Ki Kim, Shayegan Omidshafiei, Jaehong Yoon, Yue Zhang, Mohit Bansal

3. Training-free guidance in text-to-video generation via multimodal planning and structured noise initialization

[\[arxiv\]](#)[\[github\]](#)

Jialu Li*, Shoubin Yu*, **Han Lin***, Jaemin Cho, Jaehong Yoon, Mohit Bansal

PUBLICATIONS

1. Exploring MLLM-Diffusion Information Transfer with MetaCanvas

[\[arxiv\]](#)[\[project page\]](#)

Han Lin, Xichen Pan, Ziqi Huang, Ji Hou, Jialiang Wang, Weifeng Chen, Zecheng He, Felix Junfei-Xu, Junzhe Sun, Zhipeng Fan, Ali Thabet, Jaemin Cho, Mohit Bansal, Chu Wang
Advances in Neural Information Processing Systems (**NeurIPS**), 2026

2. PhyMotion: Structured 3D Motion Reward for Physics-Grounded Human Video Generation

[\[arxiv\]](#) [\[github\]](#)

Yidong Huang, Zun Wang, **Han Lin**, Dong-Ki Kim, Shayegan Omidshafiei, Jaehong Yoon, Jaemin Cho, Yue Zhang, Mohit Bansal

Advances in Neural Information Processing Systems (**NeurIPS**), 2026

3. V-Co: A Closer Look at Visual Representation Alignment via Co-Denoising

[\[arxiv\]](#) [\[project page\]](#)

Han Lin, Xichen Pan, Zun Wang, Yue Zhang, Chu Wang, Jaemin Cho, Mohit Bansal

The 19th European Conference on Computer Vision (**ECCV**), 2026

4. AnchorWeave: World-Consistent Video Generation with Retrieved Local Spatial Memories

[\[arxiv\]](#) [\[project page\]](#)

Zun Wang, **Han Lin**, Jaehong Yoon, Jaemin Cho, Yue Zhang, Mohit Bansal

The 19th European Conference on Computer Vision (**ECCV**), 2026

5. Error-Driven Scene Editing for 3D Grounding in Large Language Models

[\[arxiv\]](#) [\[github\]](#)

Yue Zhang, Zun Wang, **Han Lin**, Jialu Li, Jianing Yang, Yonatan Bitton, Idan Szpektor, Mohit Bansal

The 19th European Conference on Computer Vision (**ECCV**), 2026

6. EPiC: Efficient Video Camera Control Learning with Precise Anchor-Video Guidance

[\[arxiv\]](#) [\[project page\]](#)

Zun Wang, Jaemin Cho, Jialu Li, **Han Lin**, Jaehong Yoon, Yue Zhang, Mohit Bansal

Forty-Third International Conference on Machine Learning (**ICML**), 2026

7. A Benchmark of Expert-Level Academic Questions to Assess AI Capabilities

[\[paper\]](#) [\[project page\]](#)

Center for AI Safety, Scale AI and HLE Contributors Consortium

Nature, 2026

8. Bifrost-1: Bridging Multimodal LLMs and Diffusion Models with Patch-level CLIP Latents

[\[arxiv\]](#) [\[project page\]](#)

Han Lin, Jaemin Cho, Amir Zadeh, Chuan Li, Mohit Bansal

Advances in Neural Information Processing Systems (**NeurIPS**), 2025

9. Ctrl-Adapter: An Efficient and Versatile Framework for Adapting Diverse Controls to Any Diffusion Model

[\[arxiv\]](#) [\[project page\]](#)

Han Lin*, Jaemin Cho*, Abhay Zala, Mohit Bansal

International Conference on Learning Representations (**ICLR Oral**, top 1.82%), 2025

10. VEDiT: Latent Prediction Architecture For Procedural Video Representation Learning

[\[arxiv\]](#)

Han Lin, Tushar Nagarajan, Nicolas Ballas, Mido Assran, Mojtaba Komeili, Mohit Bansal, Koustuv Sinha

International Conference on Learning Representations (**ICLR**), 2025

11. DreamRunner: Fine-Grained Storytelling Video Generation with Retrieval-Augmented Motion Adaptation

[\[arxiv\]](#) [\[project page\]](#)

Zun Wang, Jialu Li, **Han Lin**, Jaehong Yoon, Mohit Bansal

AAAI Conference on Artificial Intelligence (**AAAI**) 2025

12. Fast Tree-Field Integrators: From Low Displacement Rank to Topological Transformers

[\[arxiv\]](#)

Krzysztof Choromanski, Arijit Sehanobish, Somnath Basu Roy Chowdhury, **Han Lin**, Avinava Dubey, Tamas Sarlos, Snigdha Chaturvedi

Advances in Neural Information Processing Systems (**NeurIPS**), 2024

13. VideoDirectorGPT: Consistent Multi-scene Video Generation via LLM-Guided Planning

[\[arxiv\]](#) [\[project page\]](#)

Han Lin, Abhay Zala, Jaemin Cho, Mohit Bansal

Conference on Language Modeling (**COLM**), 2024

14. EnvGen: Generating and Adapting Environments via LLMs for Training Embodied Agents

[\[arxiv\]](#)[\[project page\]](#)

Abhay Zala*, Jaemin Cho*, **Han Lin**, Jaehong Yoon, Mohit Bansal
Conference on Language Modeling (COLM), 2024

15. DiagrammerGPT: Generating Open-Domain, Open-Platform Diagrams via LLM Planning

[\[arxiv\]](#)[\[project page\]](#)

Abhay Zala, **Han Lin**, Jaemin Cho, Mohit Bansal
Conference on Language Modeling (COLM), 2024

16. Efficient Graph Field Integrators Meet Point Clouds

[\[arxiv\]](#)[\[github\]](#)

Krzysztof Choromanski*, Arijit Sehanobish*, **Han Lin***, Yunfan Zhao*, Eli Berger, Alvin Pan, Tetiana Parshakova, Tianyi Zhang, David Watkins, Valerii Likhoshesterov, Somnath Basu Roy Chowdhury, Avinava Dubey, Deepali Jain, Tamas Sarlos, Snigdha Chaturvedi, Adrian Weller
International Conference on Machine Learning (ICML), 2023

17. Supervised Masked Knowledge Distillation for Few-Shot Transformers

[\[arxiv\]](#)[\[github\]](#)

Han Lin*, Guangxing Han*, Jiawei Ma, Shiyuan Huang, Xudong Lin, Shih-Fu Chang
Conference on Computer Vision and Pattern Recognition (CVPR), 2023

18. Active Tactile Exploration for 3D Object Recognition

[\[arxiv\]](#)[\[project page\]](#)

Jingxi Xu*, **Han Lin***, Shuran Song, Matei Ciocarlie
IEEE International Conference on Robotics and Automation (ICRA), 2023

19. From Block-Toeplitz Matrices to Differential Equations on Graphs: Towards a General Theory for Scalable Masked Transformers

[\[arxiv\]](#)[\[github\]](#)

Krzysztof Choromanski*, **Han Lin***, Haoxian Chen*, Tianyi Zhang, Arijit Sehanobish, Valerii Likhoshesterov, Jack Parker-Holder, Tamas Sarlos, Adrian Weller, Thomas Weingarten
International Conference on Machine Learning (ICML), 2022

20. Hybrid Random Features

[\[arxiv\]](#)[\[github\]](#)

Krzysztof Choromanski*, **Han Lin***, Haoxian Chen*, Yuanzhe Ma*, Arijit Sehanobish*, Deepali Jain, Michael Ryoo, Jake Varley, Andy Zeng, Valerii Likhoshesterov, Dmitry Kalashnikov, Vikas Sindhwani, Adrian Weller
International Conference on Learning Representations (ICLR), 2022

21. Demystifying Orthogonal Monte Carlo and Beyond

[\[arxiv\]](#)[\[github\]](#)

Han Lin*, Haoxian Chen*, Tianyi Zhang, Clement Laroche, Krzysztof Choromanski
Advances in Neural Information Processing Systems (NeurIPS), 2020

* Co-First Authors, Equal Contribution

RESEARCH EXPERIENCE

Meta FAIR Lab (DReaM Team)

2026.5 - now

Research Scientist Intern

Advised by: [Marjan Ghazvininejad](#), [Xiaochuang Han](#), [Mengchen Liu](#), [Luke Zettlemoyer](#)

- Self-improvement of prompt rewriter and video generation model
- On-policy distillation (OPD) for video generation

Meta MSL (Movie Gen Team)

2025.5 - 2026.5

Research Scientist Intern

Advised by: [Chu Wang](#), [Ji Hou](#), [Jialiang Wang](#), [Weifeng Chen](#), [Zecheng He](#), [Felix Xu](#), [Junzhe Sun](#), [Zhipeng Fan](#), [Ali Thabet](#)

- Exploring VLM-Diffusion Information Transfer with MetaCanvas (under review)

Meta FAIR Lab (JEPA Team)

2024.5 - 2024.12

Research Scientist Intern

Advised by: [Koustuv Sinha](#), [Tushar Nagarajan](#), [Nicolas Ballas](#), [Mido Assran](#), [Mojtaba Komeili](#)

- VEDiT: Latent prediction architecture for procedural video representation learning (ICLR 2025)

Google Research Collaboration on Efficient Graph Networks 2022 - 2024
Advised by: [Krzysztof Choromanski](#), [Jack Parker-Holder](#), [Somnath Basu Roy Chowdhury](#), [Avinava Dubey](#), [Tamás Sarlós](#)

- Fast Tree-Field Integrators: From Low Displacement Rank to Topological Transformers (NeurIPS 2024)
- Efficient Graph Field Integrators Meet Point Clouds (ICML 2023)
- From Block-Toeplitz Matrices to Differential Equations on Graphs: Towards a General Theory for Scalable Masked Transformers (ICML 2022)

Google Research Collaboration on Random Feature Algorithms 2019 - 2022
Advised by: [Krzysztof Choromanski](#), [Deepali Jain](#), [Michael S. Ryoo](#), [Jake Varley](#), [Andy Zeng](#), [Dmitry Kalashnikov](#), [Vikas Sindhwani](#)

- Hybrid Random Features (ICLR 2022)
- Demystifying Orthogonal Monte Carlo and Beyond (NeurIPS 2020)

MURGe-Lab, UNC 2023 - now
Research Assistant, Advised by Prof. [Mohit Bansal](#)

- Text-to-video generation, multimodal learning, and LLMs

DVMM Lab, Columbia University 2022 - 2023
Research Assistant, Advised by Prof. [Shih-Fu Chang](#)

- Supervised masked knowledge distillation for few-shot Transformers (CVPR 2023)

ROAM Lab, Columbia University 2022 - 2023
Research Assistant, Advised by Prof. [Matei Ciocarlie](#) and Prof. [Shuran Song](#)

- Active tactile exploration for 3D object recognition (ICRA 2023)

Cornell, Maryland, Max Planck Pre-doctoral Research School 2022

INDUSTRY EXPERIENCE

China Merchant Securities Aug. 2020 - Aug. 2021
Option Market Making Quant Trader, Full Time

- Commodity options and futures trading and daily P&L attribution

TEACHING EXPERIENCE

COMS 4231 Analysis of Algorithms	Fall 2022
COMS 4732 Computer Vision II: Learning	Spring 2022
COMS 4721 Machine Learning for Data Science	Spring 2022
QMSS 5073 Machine Learning for Social Science	Fall 2021
IEOR 4007 Optimization Models & Methods for FE	Fall 2019
IEOR 4418 Transportation Analytics & Logistics	Spring 2019

AWARDS

CVPR 2023 Travel Award CVPR community	Jun 2023
Honorable Mention in Mathematical Context in Modeling The Consortium for Mathematics and Its Applications (COMAP)	Jan 2018
Scholarship for Excellence in Organization and Leadership Central University of Finance and Economics	Apr 2017
2nd Price for Summer & Winder Social Practice Scholarship Central University of Finance and Economics	Oct 2015

SKILLS

Python, C/C++, MATLAB, R, MySQL, L^AT_EX
PyTorch, TensorFlow, Large-Scale TPU/GPU Training

SERVICES

Workshop Organizer: Any-To-Any Multimodal Learning Workshop, CVPR, 2026
Conference Reviewer: ICML 2022-2026, NeurIPS 2022-2026, ICLR 2024-2025, CVPR 2025-2026, ICCV 2025, ECCV 2026
Conference Volunteer: Robotics: Science and Systems (RSS) 2022

REFERENCE

Mohit Bansal (mbansal@cs.unc.edu)
Department of Computer Science, UNC Chapel Hill
Relationship: PhD advisor, director of UNC MURGe-Lab.

Krzysztof Choromanski (kchoro@google.com)
Google Deepmind Robotics & Columbia University
Relationship: master's advisor and research collaborator.

Shih-Fu Chang (sc250@columbia.edu)
Columbia University
Relationship: master's research advisor, dean of Columbia Engineering.

Matei Ciocarlie (matei.ciocarlie@columbia.edu)
Columbia University
Relationship: master's research advisor, director of ROAM Lab at Columbia University.

Shuran Song (shuran@stanford.edu)
Stanford University
Relationship: master's research advisor, director of REAL Lab at Stanford University.

Koustuv Sinha (koustuvs@meta.com)
Meta
Relationship: internship advisor, research scientist at Meta FAIR Lab.

Chu Wang (wangchu@meta.com)
Meta
Relationship: internship advisor, research scientist at Meta MSL.

Marjan Ghazvininejad (ghazvini@meta.com)
Meta
Relationship: internship advisor, research scientist at Meta FAIR Lab.